

# SGN-41007 Pattern Recognition and Machine Learning

Exercise Set 5: November 25–November 29, 2019

Exercises consist of both pen&paper and computer assignments. Pen&paper questions are solved at home before exercises, while computer assignments are solved during exercise hours. The computer assignments are marked by `python` and Pen&paper questions by `pen&paper`

1. `pen&paper` Count the number of parameters in a *dense* neural network.

Consider the traditional shallow neural network architecture of Figure 1. Suppose our inputs are  $224 \times 224$  RGB bitmaps of 10 categories.

Let the network structure be the following:

- The input is  $224 \times 224 \times 3 = 150528$ -dimensional
- On the 1st layer there are 100 nodes (marked in blue)
- On the 2nd layer there are 100 nodes (marked in blue)
- On the 3rd (output) layer there are 10 nodes (marked in blue; one for each class)

Compute the number of parameters (coefficients) in the net.




Figure 1: Vanilla neural network.

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_49 (Conv2D)	(None, 64, 64, 32)	
max_pooling2d_47 (MaxPooling)	(None, 16, 16, 32)	
conv2d_50 (Conv2D)	(None, 16, 16, 32)	
max_pooling2d_48 (MaxPooling)	(None, 4, 4, 32)	
flatten_15 (Flatten)	(None, 512)	
dense_29 (Dense)	(None, 100)	
dense_30 (Dense)	(None, 2)	
Total params: 79,566		
Trainable params: 79,566		
Non-trainable params: 0		

Figure 2: A CNN model summary as reported by Keras

2. **pen&paper** Count the number of parameters in a *convolutional* neural network. Consider the Keras model defined in Listing 2. Inputs are  $64 \times 64$  color images from two categories, and all convolutional kernels are of size  $w \times h = 5 \times 5$ .
- Draw a diagram of the network.
  - Compute the number of parameters for each layer.
  - How many scalar multiplications take place on the first convolutional layer?

3. **python** *MNIST with custom network.*

Let us train two models for the MNIST dataset. Download a template code from

<http://www.cs.tut.fi/courses/SGN-41007/mnist.py>

Execute the code. The epochs should take roughly 5-10 seconds if running on the GPU. If it seems slow, try running the code directly from command line instead of Spyder. Thus; assuming the code is at C:\Temp\, you do the following:

```
>> c:
>> cd \Temp
>> python mnist.py
```

4. **python** *MNIST with pretrained network.*

Instead of using the custom ConvNet of question 4, initialize a mobilenet and attach 2 dense layers at the end. The mobilenet should have parameters `include_top = False`, `alpha = 0.25`, and `input_shape = (128, 128, 3)`. Check `model.summary()`.

Add dense layers after the convolutional pipeline such that `model.summary()` reports the following (top of listing omitted):

<code>conv_pw_13_bn</code>	<code>(BatchNormaliz (None, 4, 4, 256)</code>	<code>1024</code>
<code>conv_pw_13_relu</code>	<code>(ReLU) (None, 4, 4, 256)</code>	<code>0</code>
<code>flatten</code>	<code>(Flatten) (None, 4096)</code>	<code>0</code>
<code>dense</code>	<code>(Dense) (None, 100)</code>	<code>409700</code>
<code>dense_1</code>	<code>(Dense) (None, 10)</code>	<code>1010</code>
=====		
<code>Total params: 629,254</code>		
<code>Trainable params: 623,782</code>		
<code>Non-trainable params: 5,472</code>		

5. **python** *Train the model.*

After designing the model, compile and train it. You will also need to resize your inputs to match the input size of the net (128 x 128). What accuracy do you get?